

My SQL - Assignment 2

Assignment 1: Write a SELECT query to retrieve all columns from a 'customers' table, and modify it to return only the customer name and email address for customers in a specific city.

Task 1: Retrieve All Columns from the 'customers' Table

```
mysql> select * from customers;
```

CusId	CusName	email	city	phoneNo
10	Sam	sam@ex.com	Hyderabad	987654321
20	Tom	tom@ex.com	Chennai	876543292
30	Ram	ram@ex.com	delhi	976543289
40	Bob	bob@ex.com	Hyderabad	899765534
50	Suz	suz@ex.com	Banglore	677889578
60	Sasi	sasi@ex.com	Hyderabad	777654321
70	Aman	aman@ex.com	Chennai	886543222
80	Ravi	ravi@ex.com	Hyderabad	976543266
90	Roy	roy@ex.com	Hyderabad	699764434
100	Jake	jake@ex.com	Banglore	777889778

```
10 rows in set (0.00 sec)
```

Task 2: Retrieve Customer Name and Email Address for Customers in a Specific City

```
mysql> select CusName, email from customers
-> where city = 'Hyderabad';
```

CusName	email
Sam	sam@ex.com
Bob	bob@ex.com
Sasi	sasi@ex.com
Ravi	ravi@ex.com
Roy	roy@ex.com

```
5 rows in set (0.00 sec)
```

Assignment 2: Craft a query using an INNER JOIN to combine 'orders' and 'customers' tables for customers in a specified region, and a LEFT JOIN to display all customers including those without orders.

Query 1:

INNER JOIN for Customers in a Specified Region :

```
mysql> SELECT
->     customers.CusId,
->     customers.CusName,
->     customers.email,
->     orders.OrderId,
->     orders.OrderDate,
->     orders.Amount
-> FROM
->     customers
-> INNER JOIN
->     orders ON customers.CusId = orders.CusId
-> WHERE
->     customers.region = 'South';
```

CusId	CusName	email	OrderId	OrderDate	Amount
10	Sam	sam@ex.com	1	2024-05-01	150.00
10	Sam	sam@ex.com	3	2024-05-03	100.00
10	Sam	sam@ex.com	22	2024-05-22	1150.00
20	Tom	tom@ex.com	2	2024-05-02	200.00
40	Bob	bob@ex.com	5	2024-05-05	300.00
40	Bob	bob@ex.com	24	2024-05-24	1250.00
50	Suz	suz@ex.com	6	2024-05-06	350.00
60	Sasi	sasi@ex.com	7	2024-05-07	400.00
60	Sasi	sasi@ex.com	25	2024-05-25	1300.00
70	Aman	aman@ex.com	8	2024-05-08	450.00
70	Aman	aman@ex.com	26	2024-05-26	1350.00
80	Ravi	ravi@ex.com	9	2024-05-09	500.00
80	Ravi	ravi@ex.com	27	2024-05-27	1400.00
90	Roy	roy@ex.com	10	2024-05-10	550.00
100	Jake	jake@ex.com	11	2024-05-11	600.00
100	Jake	jake@ex.com	28	2024-05-28	1450.00

16 rows in set (0.00 sec)

Query 2:

LEFT JOIN to Display All Customers Including Those Without Orders :

```
mysql> SELECT
->     customers.CusId,
->     customers.CusName,
->     customers.email,
->     orders.OrderId,
->     orders.OrderDate,
->     orders.Amount
-> FROM
->     customers
-> LEFT JOIN
->     orders ON customers.CusId = orders.CusId;
```

CusId	CusName	email	OrderId	OrderDate	Amount
10	Sam	sam@ex.com	1	2024-05-01	150.00
10	Sam	sam@ex.com	3	2024-05-03	100.00
10	Sam	sam@ex.com	22	2024-05-22	1150.00
20	Tom	tom@ex.com	2	2024-05-02	200.00
30	Ram	ram@ex.com	4	2024-05-04	250.00
30	Ram	ram@ex.com	23	2024-05-23	1200.00
40	Bob	bob@ex.com	5	2024-05-05	300.00
40	Bob	bob@ex.com	24	2024-05-24	1250.00
50	Suz	suz@ex.com	6	2024-05-06	350.00
60	Sasi	sasi@ex.com	7	2024-05-07	400.00
60	Sasi	sasi@ex.com	25	2024-05-25	1300.00
70	Aman	aman@ex.com	8	2024-05-08	450.00
70	Aman	aman@ex.com	26	2024-05-26	1350.00
80	Ravi	ravi@ex.com	9	2024-05-09	500.00
80	Ravi	ravi@ex.com	27	2024-05-27	1400.00
90	Roy	roy@ex.com	10	2024-05-10	550.00
100	Jake	jake@ex.com	11	2024-05-11	600.00
100	Jake	jake@ex.com	28	2024-05-28	1450.00
110	Sruthi	sruthi@ex.com	NULL	NULL	NULL
120	Amar	amar@ex.com	13	2024-05-13	700.00
120	Amar	amar@ex.com	30	2024-05-30	1550.00
130	Ravan	ravan@ex.com	14	2024-05-14	750.00
130	Ravan	ravan@ex.com	31	2024-05-31	1600.00
140	Krish	krish@ex.com	NULL	NULL	NULL
150	Adam	adam@ex.com	16	2024-05-16	850.00
160	Vasu	vasu@ex.com	17	2024-05-17	900.00
170	Smith	smith@ex.com	18	2024-05-18	950.00
180	jay	jay@ex.com	NULL	NULL	NULL
190	Arjun	arjun@ex.com	NULL	NULL	NULL
200	Charan	charan@ex.com	21	2024-05-21	1100.00

30 rows in set (0.00 sec)

Assignment 3: Utilize a subquery to find customers who have placed orders above the average order value, and write a UNION query to combine two SELECT statements with the same number of columns.

Query 1:

Subquery to Find Customers with Orders Above the Average Order Value

```
mysql> SELECT
->     customers.CusId,
->     customers.CusName,
->     customers.email,
->     orders.OrderId,
->     orders.OrderDate,
->     orders.Amount
-> FROM
->     customers
-> INNER JOIN
->     orders ON customers.CusId = orders.CusId
-> WHERE
->     orders.amount > (SELECT AVG(amount) FROM orders);
```

CusId	CusName	email	OrderId	OrderDate	Amount
150	Adam	adam@ex.com	16	2024-05-16	850.00
160	Vasu	vasu@ex.com	17	2024-05-17	900.00
170	Smith	smith@ex.com	18	2024-05-18	950.00
200	Charan	charan@ex.com	21	2024-05-21	1100.00
10	Sam	sam@ex.com	22	2024-05-22	1150.00
30	Ram	ram@ex.com	23	2024-05-23	1200.00
40	Bob	bob@ex.com	24	2024-05-24	1250.00
60	Sasi	sasi@ex.com	25	2024-05-25	1300.00
70	Aman	aman@ex.com	26	2024-05-26	1350.00
80	Ravi	ravi@ex.com	27	2024-05-27	1400.00
100	Jake	jake@ex.com	28	2024-05-28	1450.00
120	Amar	amar@ex.com	30	2024-05-30	1550.00
130	Ravan	ravan@ex.com	31	2024-05-31	1600.00

13 rows in set (0.00 sec)

Query 2:

UNION Query to Combine Two SELECT Statements

Let's assume we want to combine:

- A list of customers from the "South" region.
- A list of customers who have placed an order amount greater than \$1000.

```
mysql> SELECT
->     CusId,
->     CusName,
->     email,
->     city,
->     phoneNo,
->     region
-> FROM
->     customers
-> WHERE
->     region = 'South'
-> UNION
-> SELECT
->     c.CusId,
->     c.CusName,
->     c.email,
->     c.city,
->     c.phoneNo,
->     c.region
-> FROM
->     customers c
-> INNER JOIN
->     orders o ON c.CusId = o.CusId
-> WHERE
->     o.amount > 1000;
```

CusId	CusName	email	city	phoneNo	region
10	Sam	sam@ex.com	Hyderabad	987654321	South
20	Tom	tom@ex.com	Chennai	876543292	South
40	Bob	bob@ex.com	Hyderabad	899765534	South
50	Suz	suz@ex.com	Banglore	677889578	South
60	Sasi	sasi@ex.com	Hyderabad	777654321	South
70	Aman	aman@ex.com	Chennai	886543222	South
80	Ravi	ravi@ex.com	Hyderabad	976543266	South
90	Roy	roy@ex.com	Hyderabad	699764434	South
100	Jake	jake@ex.com	Banglore	777889778	South
200	Charan	charan@ex.com	Jaipur	700889778	West
30	Ram	ram@ex.com	delhi	976543289	North
120	Amar	amar@ex.com	kolkata	686543211	East
130	Ravan	ravan@ex.com	Jaipur	676543299	West

13 rows in set (0.00 sec)

Assignment 4: Compose SQL statements to BEGIN a transaction, INSERT a new record into the 'orders' table, COMMIT the transaction, then UPDATE the 'products' table, and ROLLBACK the transaction.

Step 1: BEGIN a transaction and INSERT a new record into the 'orders' table, then COMMIT the transaction

```
mysql> START TRANSACTION;
Query OK, 0 rows affected (0.00 sec)

mysql> INSERT INTO Orders(OrderId,CusId,OrderDate,Amount)
-> VALUES (32,10,'2024-06-01',500.00);
Query OK, 1 row affected (0.00 sec)

mysql> COMMIT;
Query OK, 0 rows affected (0.00 sec)

mysql> SELECT * FROM Orders;
```

OrderId	CusId	OrderDate	Amount
1	10	2024-05-01	150.00
2	20	2024-05-02	200.00
3	10	2024-05-03	100.00
4	30	2024-05-04	250.00
5	40	2024-05-05	300.00
6	50	2024-05-06	350.00
7	60	2024-05-07	400.00
8	70	2024-05-08	450.00
9	80	2024-05-09	500.00
10	90	2024-05-10	550.00
11	100	2024-05-11	600.00
13	120	2024-05-13	700.00
14	130	2024-05-14	750.00
16	150	2024-05-16	850.00
17	160	2024-05-17	900.00
18	170	2024-05-18	950.00
21	200	2024-05-21	1100.00
22	10	2024-05-22	1150.00
23	30	2024-05-23	1200.00
24	40	2024-05-24	1250.00
25	60	2024-05-25	1300.00
26	70	2024-05-26	1350.00
27	80	2024-05-27	1400.00
28	100	2024-05-28	1450.00
30	120	2024-05-30	1550.00
31	130	2024-05-31	1600.00
32	10	2024-06-01	500.00

```
27 rows in set (0.00 sec)
```

Step 2: BEGIN another transaction, **UPDATE** the 'products' table, and **ROLLBACK** the transaction

```
mysql> START TRANSACTION;
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> UPDATE product SET stock = stock-10
      -> where ProductId=1;
Query OK, 1 row affected (0.00 sec)
Rows matched: 1  Changed: 1  Warnings: 0
```

```
mysql> Select * from Product;
```

ProductId	ProductName	Price	Stock
1	Mobile Phone	15000.00	40
2	Laptop	50000.00	30
3	Headphones	2000.00	100
4	Smart Watch	10000.00	40
5	Tablet	25000.00	20
6	Camera	30000.00	15
7	Bluetooth Speaker	5000.00	60
8	Television	40000.00	10
9	Gaming Console	35000.00	8
10	Refrigerator	45000.00	12

```
10 rows in set (0.00 sec)
```

```
mysql> ROLLBACK;
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> Select * from Product;
```

ProductId	ProductName	Price	Stock
1	Mobile Phone	15000.00	50
2	Laptop	50000.00	30
3	Headphones	2000.00	100
4	Smart Watch	10000.00	40
5	Tablet	25000.00	20
6	Camera	30000.00	15
7	Bluetooth Speaker	5000.00	60
8	Television	40000.00	10
9	Gaming Console	35000.00	8
10	Refrigerator	45000.00	12

```
10 rows in set (0.00 sec)
```

Assignment 5: Begin a transaction, perform a series of INSERTs into 'orders', setting a SAVEPOINT after each, rollback to the second SAVEPOINT, and COMMIT the overall transaction.

Steps :

- Begins a transaction using the **START**; statement.
- Inserts three orders into the 'orders' table, setting a **SAVEPOINT** after each **INSERT** operation.
- Rolls back to the second **SAVEPOINT** (savepoint2).
- Finally, commits the overall transaction using the **COMMIT**; statement.

```
mysql> START TRANSACTION;
Query OK, 0 rows affected (0.01 sec)

mysql> INSERT INTO Orders (OrderId,CusId,OrderDate,Amount)
-> VALUES (33,10,'2024-06-01',500.00);
Query OK, 1 row affected (0.01 sec)

mysql> SAVEPOINT savepoint1;
Query OK, 0 rows affected (0.00 sec)

mysql> INSERT INTO Orders (OrderId,CusId,OrderDate,Amount)
-> VALUES (34,20,'2024-07-01',700.00);
Query OK, 1 row affected (0.00 sec)

mysql> SAVEPOINT savepoint2;
Query OK, 0 rows affected (0.00 sec)

mysql> INSERT INTO Orders (OrderId,CusId,OrderDate,Amount)
-> VALUES (35,30,'2024-06-05',900.00);
Query OK, 1 row affected (0.00 sec)

mysql> SAVEPOINT savepoint3;
Query OK, 0 rows affected (0.00 sec)

mysql> ROLLBACK TO SAVEPOINT savepoint2;
Query OK, 0 rows affected (0.00 sec)

mysql> COMMIT;
Query OK, 0 rows affected (0.00 sec)
```

In the picture below, you can see that the first two insertions are done, but the third insertion is rolled back. That's why you can't see the OrderId with 35.


```
mysql> SELECT OrderId,CusId,OrderDate,Amount FROM Orders;
```

OrderId	CusId	OrderDate	Amount
1	10	2024-05-01	150.00
2	20	2024-05-02	200.00
3	10	2024-05-03	100.00
4	30	2024-05-04	250.00
5	40	2024-05-05	300.00
6	50	2024-05-06	350.00
7	60	2024-05-07	400.00
8	70	2024-05-08	450.00
9	80	2024-05-09	500.00
10	90	2024-05-10	550.00
11	100	2024-05-11	600.00
13	120	2024-05-13	700.00
14	130	2024-05-14	750.00
16	150	2024-05-16	850.00
17	160	2024-05-17	900.00
18	170	2024-05-18	950.00
21	200	2024-05-21	1100.00
22	10	2024-05-22	1150.00
23	30	2024-05-23	1200.00
24	40	2024-05-24	1250.00
25	60	2024-05-25	1300.00
26	70	2024-05-26	1350.00
27	80	2024-05-27	1400.00
28	100	2024-05-28	1450.00
30	120	2024-05-30	1550.00
31	130	2024-05-31	1600.00
32	10	2024-06-01	500.00
33	10	2024-06-01	500.00
34	20	2024-07-01	700.00

```
29 rows in set (0.00 sec)
```

Assignment 6: Draft a brief report on the use of transaction logs for data recovery and create a hypothetical scenario where a transaction log is instrumental in data recovery after an unexpected shutdown.

Report on the Use of Transaction Logs for Data Recovery :

Transaction logs play a vital role in ensuring data integrity and facilitating recovery in the event of system failures or unexpected shutdowns. These logs record all changes made to the database during transactions, providing a detailed history of data modifications.

Scenario: Imagine a scenario where an online retail company experiences a sudden power outage during a busy sales period. As a result, the database server abruptly shuts down, leading to potential data corruption and loss. However, due to the presence of transaction logs, the company can recover the lost data efficiently.

Utilization of Transaction Logs:

1. **Recovery Point:** Transaction logs serve as a recovery point, allowing the database to be restored to a specific point in time before the failure occurred.
2. **Redo and Undo Operations:** The transaction logs contain both redo and undo information. Redo logs help reapply committed transactions that were not yet written to disk before the failure. Undo logs assist in rolling back uncommitted or partially committed transactions to maintain data consistency.
3. **Consistency Check:** Before applying the redo and undo operations, the integrity of the transaction logs is verified to ensure their accuracy and completeness.

Conclusion: In conclusion, transaction logs are indispensable for data recovery in the event of unexpected system failures. By capturing all database modifications in real-time, these logs provide a reliable mechanism for restoring data integrity and minimizing downtime, thereby ensuring business continuity.